



Build Cache Deep Dive

Improving the Developer Experience with Faster Feedback Cycles

About the trainer

Prerequisites

- Skills
 - Good understanding of Java language
 - Basic understanding of Gradle concepts
- Tools
 - Java 8
 - Latest Gradle version

Training content

- Build cache in context
- Understand the benefits of using the Gradle build cache
- Use and configure the build cache
- Tune build logic for maximum cacheability
- Maximize the benefits with Gradle Enterprise

Training material

- Gradle Enterprise training instance
@ [https //enterprise-training.gradle.com](https://enterprise-training.gradle.com)
- Zip with hands-on labs and slides
@ [https //enterprise-training.gradle.com/build-cache-deep-dive](https://enterprise-training.gradle.com/build-cache-deep-dive)

Build scans

- Gathers details about build
- Generated and published with `--scan`
- Captures IP address and host name
- Published information will be publicly-available
- Can be deleted manually with minus icon in toolbar at the top

Performance is key

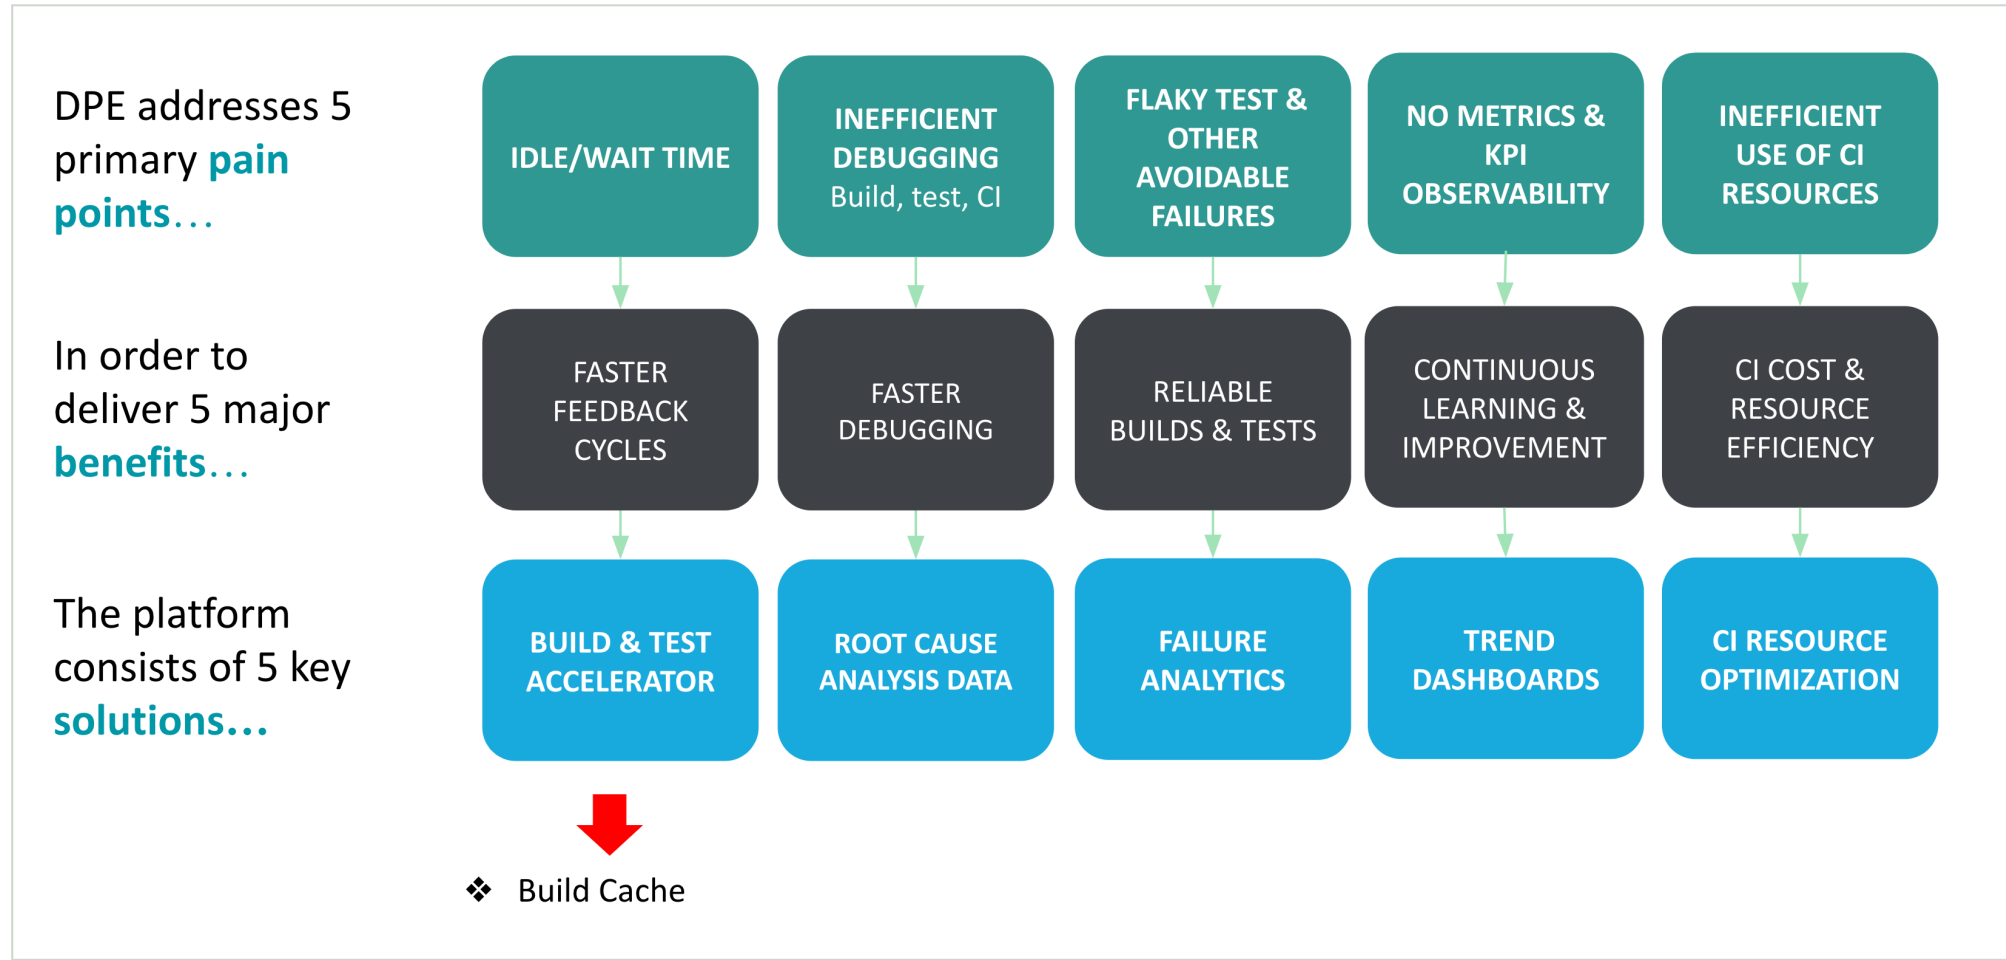
- Faster build times lead to faster feedback
- Faster feedback leads to better developer productivity
- Better developer productivity ships features quicker
- See blog post [Quantifying the costs of builds](#)

Build cache in context

Build Cache is a key enabling technology for a new software development discipline called Developer Productivity Engineering (DPE). DPE uses data analytics and acceleration technologies to speed up software development processes—from builds to testing to CI—and make them more efficient.

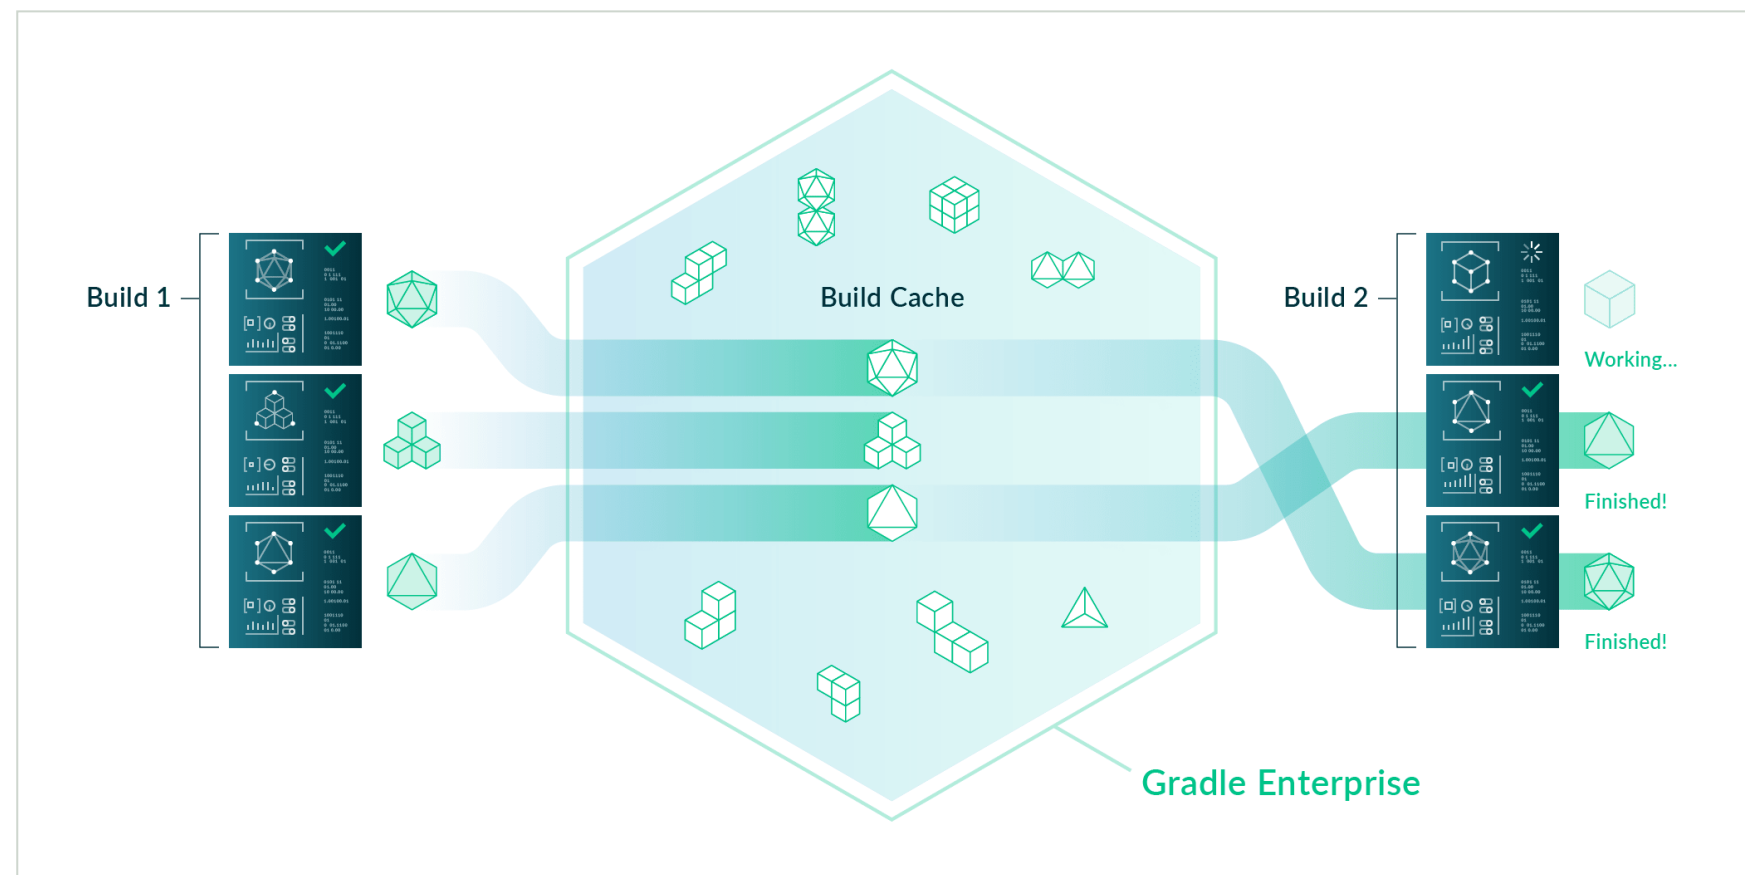
- Faster feedback cycles
- More reliable and actionable data
- A highly satisfying developer experience

Developer productivity engineering solution framework



Approaches for build avoidance


- **Incremental build** build avoidance in same workspace
- **Build cache** share build results across multiple workspaces



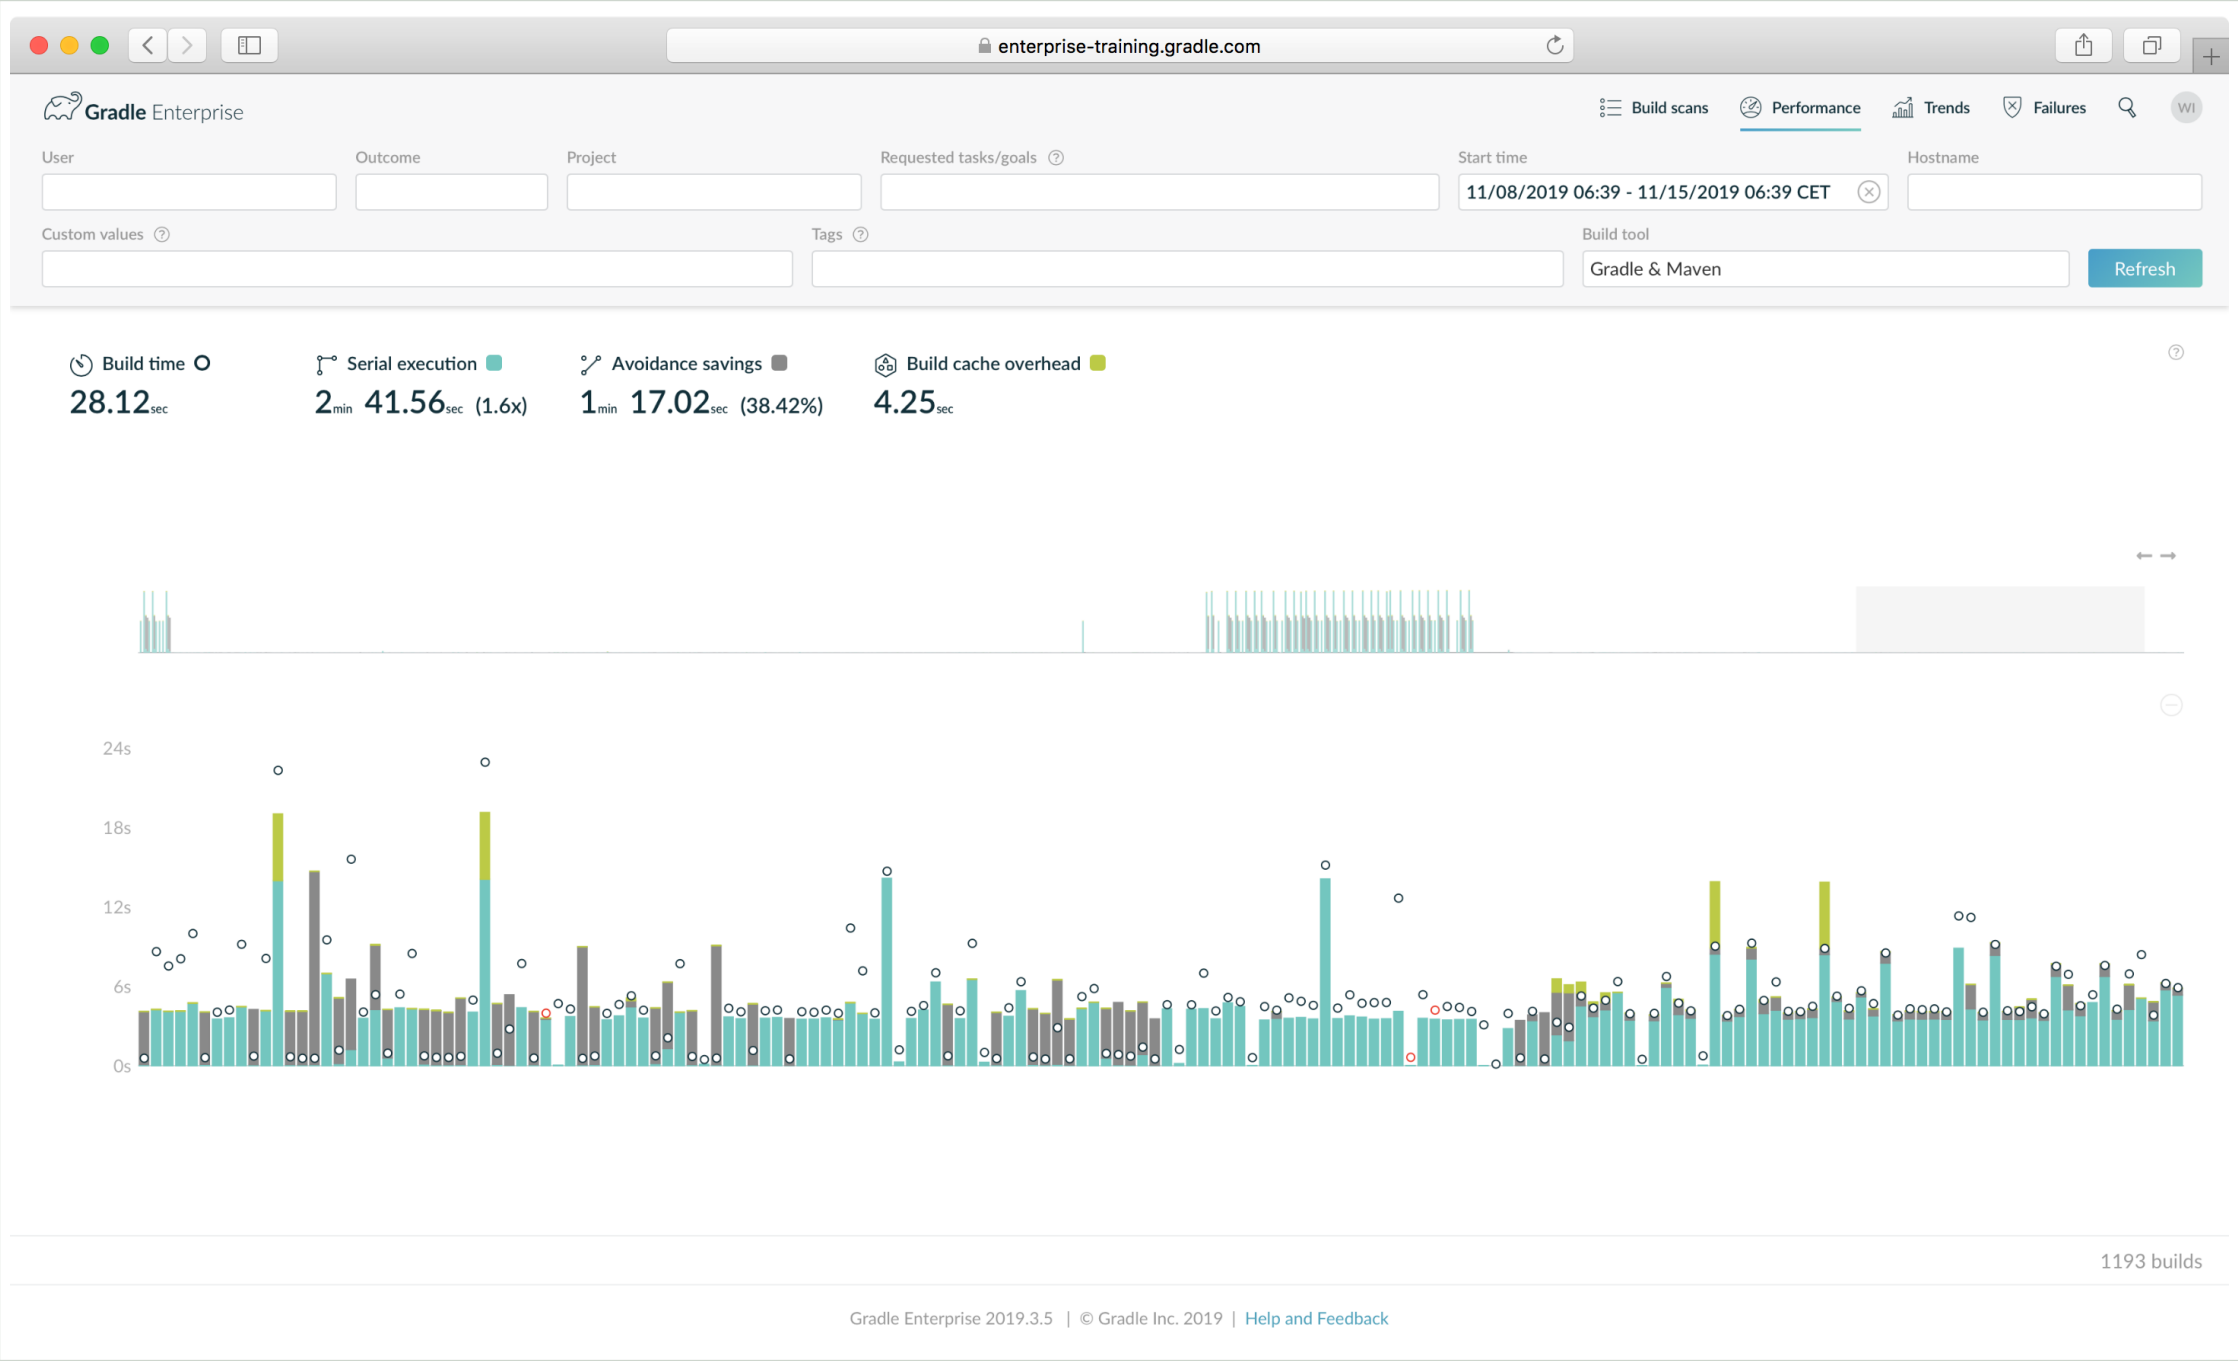
Common use cases

- Speed up developers' builds when switching branches
- Share results between CI builds
- Accelerate developer builds by reusing CI results
- Driving positive developer behavior change with faster build & test feedback cycles

Spring Boot build time for compile & unit tests



Visualized savings in Gradle Enterprise




Recap incremental builds

- Important feature from the beginning
- Optimized for single developer running the build
- Underlying mechanism for the build cache

How does it work?

- Task needs to define inputs and outputs
- Hashes of inputs and outputs are stored on disk
- Actions are only executed if inputs and/or outputs have changed



Execution marker in console

- Gradle marks task **UP-TO-DATE**
- Build summary indicates high-level status

```
$ gradle compileJava --console=verbose  
:compileJava UP-TO-DATE
```

```
BUILD SUCCESSFUL in 0s  
1 actionable task: 1 up-to-date
```


Declaring inputs and outputs with annotations

Generate.groovy

```
class Generate extends DefaultTask {  
    @Input  
    int fileCount = 10  
  
    @OutputDirectory  
    File generatedFileDir = project.file("${project.buildDir}/generated")  
  
    @TaskAction  
    void perform() {  
        for (int i=0; i<fileCount; i++) {  
            new File(generatedFileDir, "${i}.txt").text = i  
        }  
    }  
}
```

Assign annotations to task properties or getter methods for all of your custom task implementations.

Declaring inputs and outputs with run me API

build.gradle

```
generate {  
    inputs.property 'fileCount', 10  
    outputs.dir project.file("${project.buildDir}/generated")  
}
```

Use run me task API (see **TaskInputs** and **TaskOutputs**) if task source code cannot be changed easily.

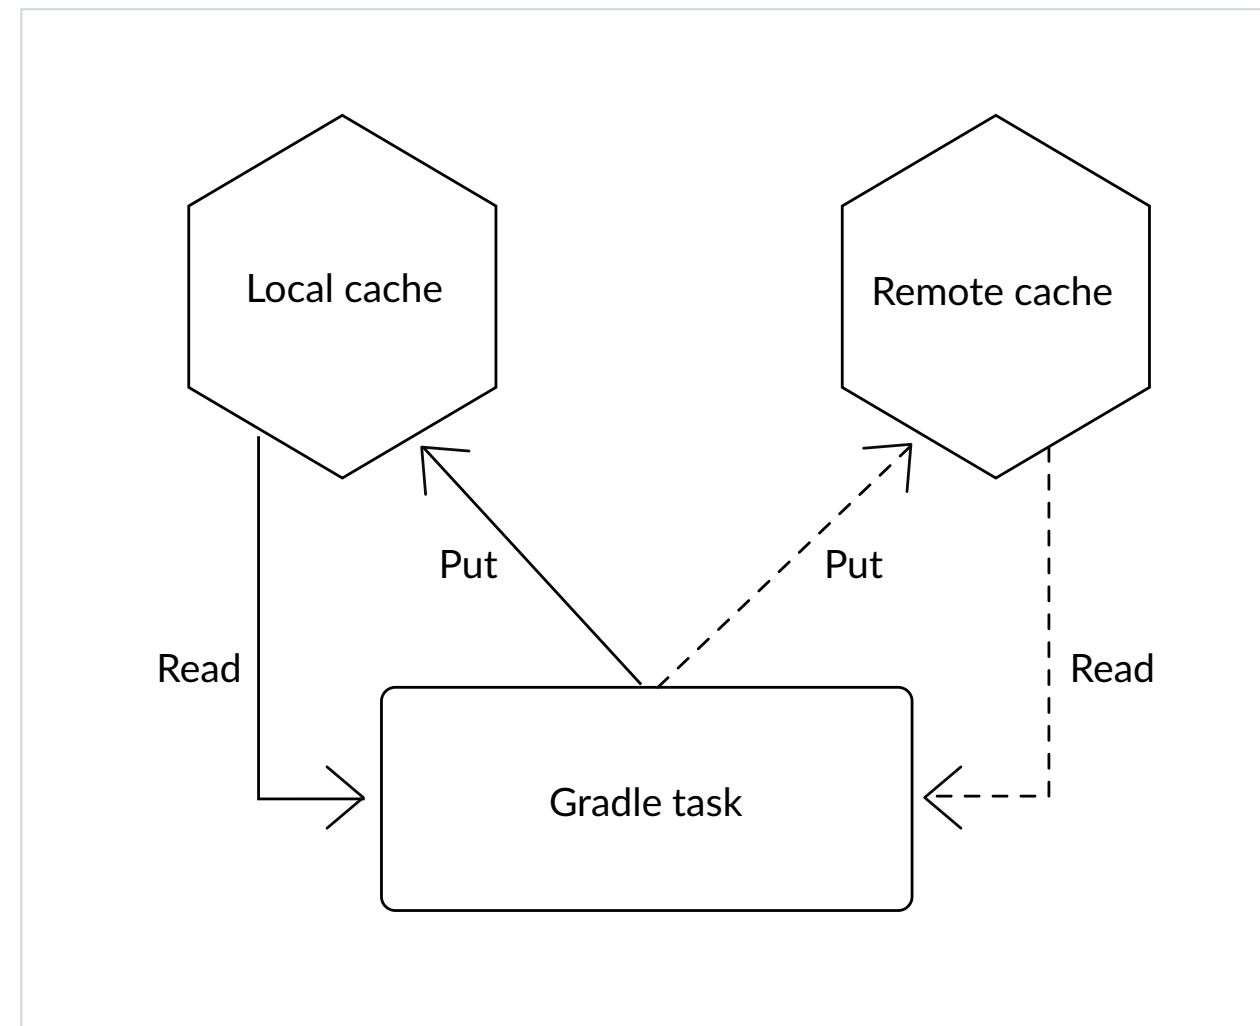
What are the limitations?

- Only uses the result of the previous execution
- Restricted to execution on single machine
- Cache is not shared among team members

What is the build cache?

- Reuse build outputs of *any* previous execution
- Reuse build outputs even if run with `clean` task
- Uniquely identifies outputs of tasks by inputs
- First stable version with Gradle 4.0

Different types of build caches



Local build cache

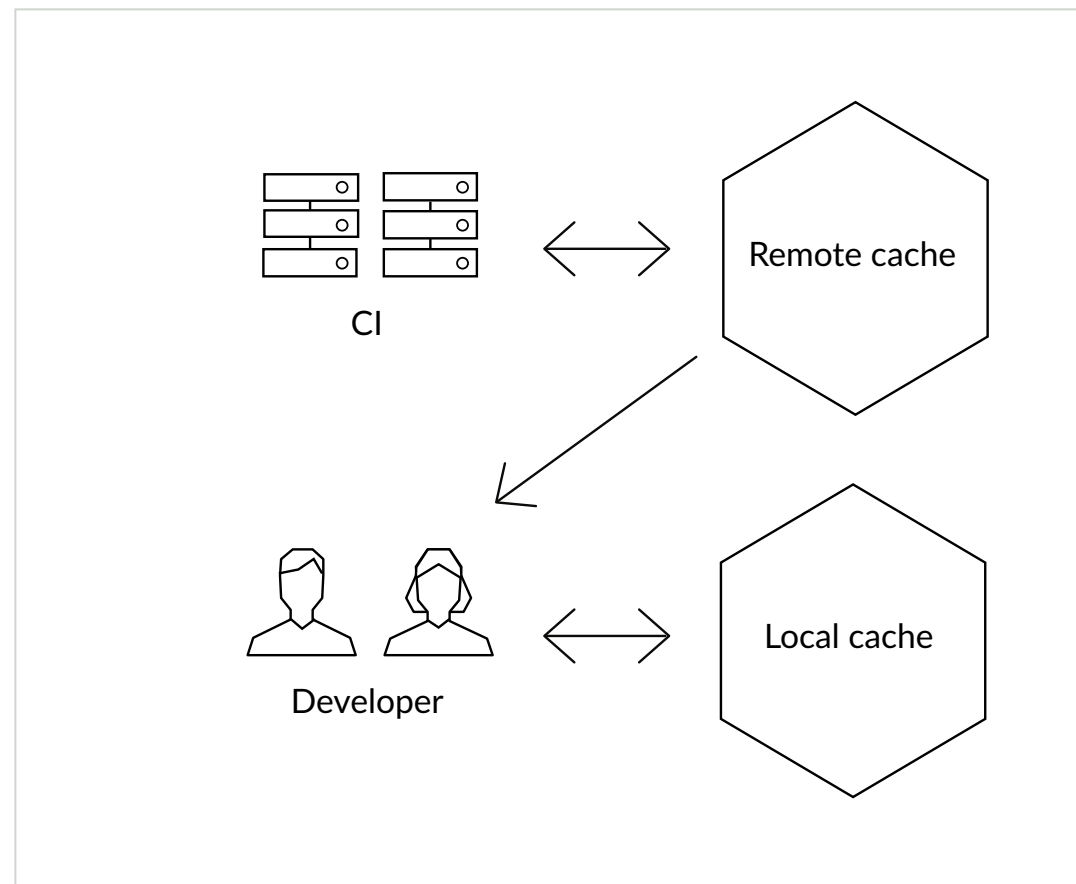
- Uses cache in directory on local machine
- Speeds up development for single developer or build agent
- Reuses build results when switching branches locally
- Particularly useful for Android variants

Remote build cache

- Shared among different machines
- Speeds up development for the whole team
- Reuses build results among CI agents/jobs and individual developers

Recommended sharing strategy

- only push to the shared cache from CI
- avoid sharing from developer machines



Influencing factors

- Architecture of code
- Nature of change
- Are tasks cacheable?
- Do outputs change with every build?

Using the build cache

Enabling the cache

- This build invocation only `--build-cache` command line option
- All build invocations `org.gradle.caching=true` in `gradle.properties`
- The `buildSrc` project needs to be **explicitly enabled for cacheability**

```
gradle --build-cache clean assemble
```

Lab 01

Using the local build cache

Configuring the local build cache

se ngs.gradle

```
buildCache {  
    local(DirectoryBuildCache) {  
        directory = new File(rootDir, 'build-cache')  
        removeUnusedEntriesAfterDays = 30  
    }  
}
```

Domain class for configuring local cache **DirectoryBuildCache**

Configuring the remote build cache

se ngs.gradle

```
buildCache {  
    remote(HttpBuildCache) {  
        url = 'http://example.com:8123/cache/'  
        credentials {  
            username = 'build-cache-user'  
            password = 'some-complicated-password'  
        }  
    }  
}
```

Domain class for configuring remote cache **HttpBuildCache**

Conditional cache configuration

see ngs.gradle

```
def ciServer = System.getenv().containsKey('CI')

buildCache {
    local {
        enabled = !ciServer
    }
    remote(HttpBuildCache) {
        url = 'https://example.com:8123/cache/'
        push = ciServer
    }
}
```

Standardizing build cache configuration on

init.gradle

```
def ciServer = System.getenv().containsKey('CI')


gradle.settingsEvaluated { settings ->
    settings.buildCache {
        local {
            enabled = !ciServer
        }
        remote(HttpBuildCache) {
            url = 'https://example.com:8123/cache/'
            push = ciServer
        }
    }
}
```


Computing the build cache key

- The task implementation
- The task action implementations
- The names of the output properties
- The names and values of task inputs

Build cache operations

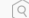


- Hit
- Miss
- Store
- Packing
- Unpacking

 Build Scan

t

✓ gradle clean largeMonolithicJavaProject fullPerfor...

Mar 9, 2018 12:17:10 PM MST



Summary

Console log

Timeline

Performance

Tests

Projects

Dependencies

Plugins

Custom values

Switches

Infrastructure

Build

Configuration

Dependency resolution

Task execution

Build cache

Network activity

Settings and suggestions

Tasks whose outputs were requested from cache

161

Hit

161 (100%)

Local

0

Remote

161 (100%)

Miss

0

Tasks whose outputs were stored to cache

0

Local cache

(disabled)

Remote cache (HTTP)

Push

enabled

Configuration

Authenticated

true

AllowUntrustedServer

false

URL

<https://e.grdev.net/cache/>

Operations

Hit >

161 5.061s 59.27 MB 11.7 MB/s

Miss

0

Store

0

Packing and unpacking ⓘ

Pack

0

Unpack >

161 7.762s 59.27 MB

Lab 02

Using the remote build cache

What makes a task cacheable ?

- Task needs to define inputs and outputs
- Task type implementation needs to declare `@CacheableTask` annotation
- `@CacheableTask` is not inherited by subclasses
- Custom task types have to opt into cacheability

Cacheability influencing factors

- Declared inputs and outputs
- Repeatable output
- Relocatability vs. absolute paths

Built-in cacheable tasks

- Some but not all built-in Gradle tasks are cacheable
- Tasks involving copy operations are usually not cacheable

Enabling cacheability by annotation

Generate.groovy

```
@CacheableTask
class Generate extends DefaultTask {
    @Input
    int fileCount = 10

    @OutputDirectory
    File generatedFileDir = project.file("${project.buildDir}/generated")

    @TaskAction
    void perform() {
        for (int i=0; i<fileCount; i++) {
            new File(generatedFileDir, "${i}.txt").text = i
        }
    }
}
```

Only applicable to custom task implementations!

Enabling cacheability by runtime API

build.gradle

```
generateCode {  
    outputs.cacheIf {  
        // return boolean expression  
    }  
}
```

Ad-hoc tasks or tasks from plugins can determine cacheability via **`TaskOutputs.cacheIf(Spec)`**.

Disabling cacheability by runtime API

build.gradle

```
generateCode {  
    outputs.doNotCacheIf('Actions produce volatile results') {  
        true  
    }  
}
```

Disabling the cache for a task with `TaskOutputs.doNotCacheIf(String, Spec)` requires providing a reason.

Lab 03

Equipping tasks with caching capabilities

Troubleshooting the build cache

Possible approaches

- Low-level troubleshooting
 - Identify task outcome with `--console=verbose`
 - Retrieve cache key information by changing the log level
 - Compare cache keys and root causes
- Visual and convenient troubleshooting
 - Create a build scan
 - Use GE deep insight features

Info log level console information

```
$ gradle helloWorld --build-cache -i  
  
> Task :helloWorld UP-TO-DATE  
Build cache key for task ':helloWorld' is 16f4fbc007345a854d49302279d1
```

Info log level displays cache key generated for each task.

Debug log level console information

```
$ gradle helloWorld --build-cache -Dorg.gradle.caching.debug=true  
  
> Task :helloWorld UP-TO-DATE  
Appending taskClass to build cache key: HelloWorld_Decorated  
Appending classLoaderHash to build cache key: 575dae0f1414d5dfd4ef14b6  
Appending actionType to build cache key: HelloWorld_Decorated  
Appending actionClassLoaderHash to build cache key: 575dae0f1414d5dfd4  
Appending inputPropertyHash for 'message' to build cache key: f81fd65e  
Appending outputPropertyName to build cache key: outputFile  
Build cache key for task ':helloWorld' is 16f4fbc007345a854d49302279d1
```

Debug log level displays more detailed information.

Using build scans

- Task input comparison
- Task details (cache key, cacheability reason)
- Determining origin build of cache output
- Performance breakdown

Requirements for cacheable tasks

Repeatable task outputs

- Same inputs should produce the same outputs
- Byte-for-byte equivalent or semantically equivalent (with normalization)

Stable task inputs

- Inputs need to be stable over time
- Potential source of volatility
 - Timestamps
 - Absolute file paths
 - Non-deterministic ordering

Path sensitivity

- File paths for input properties are absolute by default
- Shared build results between machines requires exact same path
- Controllable via annotation `@PathSensitive`

```
@PathSensitive(PathSensitivity.RELATIVE)
@InputFiles
public FileTree getSources() {
    // ...
}
```

Input normalization

- Task inputs between two executions are compared to determine cacheability
- Controllable via annotations `@Classpath` and `@CompileClasspath`
- Example For compile classpath Gradle extracts ABI signature from the classes on the classpath
- Configurable to ignore volatile files via `Project.normalization(Action)`

build.gradle

```
normalization {  
    runtimeClasspath {  
        ignore 'build-info.properties'  
    }  
}
```

Handling cases affecting cache correctness

Overlapping outputs

- Two or more tasks write to the same directory
- Difficult for Gradle to determine which output belongs to which task
- Build scan renders reason for this case

External inputs like system properties

- System properties often use absolute path
- Use relative path to fix

File encoding

- Java tools use the system file encoding when no specific encoding specified
- Can cause incorrect builds
- Always set the file system encoding to avoid issues

Line endings

- Important when build cache is shared across different OSes
- Set `autocrlf=false` if Git is used

Symlinks

- Symlinks are not stored in build cache
- Uses actual file contents of the destination of the link
- Some OSes (e.g. Windows) do not support symlinks
- Tasks will not be cacheable across different OSes

Java versions

- Gradle tracks only the major version of Java as input
- Usually applicable to compilation and test tasks
- Vendor and the minor version may influence the bytecode
- Suggested to add vendor as an input to the corresponding tasks

Lab 04

Handling cache misses

Getting started with the build cache

Recommended approach

- Equip tasks with inputs and outputs
- Use local build cache
- Set up remote build cache
- Roll out usage to team
- Use Gradle Enterprise for cache monitoring and optimization
- Report from the field [Tableau using Gradle Enterprise](#)

Installing the remote build cache

- Build cache node available as Docker image or JAR file
- Docker image freely-available from Docker Hub
- Requires Docker installation on host machine

PUBLIC REPOSITORY

gradle/build-cache-node ☆

Last pushed: 25 days ago

Repo Info Tags

Short Description

A remote Gradle build cache, capable of connecting to Gradle Enterprise.

Full Description

A [Gradle](#) build cache node operates as a remote Gradle build cache, and can connect with [Gradle Enterprise](#) for centralized management. The cache node can also be used without a Gradle Enterprise installation with restricted functionality.

For more information on installing and operating a build cache node, please see the [Gradle Enterprise Admin Manual](#).


For more information on Gradle build caching, please see the [Gradle User Guide](#).

For help and assistance, please use [discuss.gradle.org](#).

Docker Pull Command

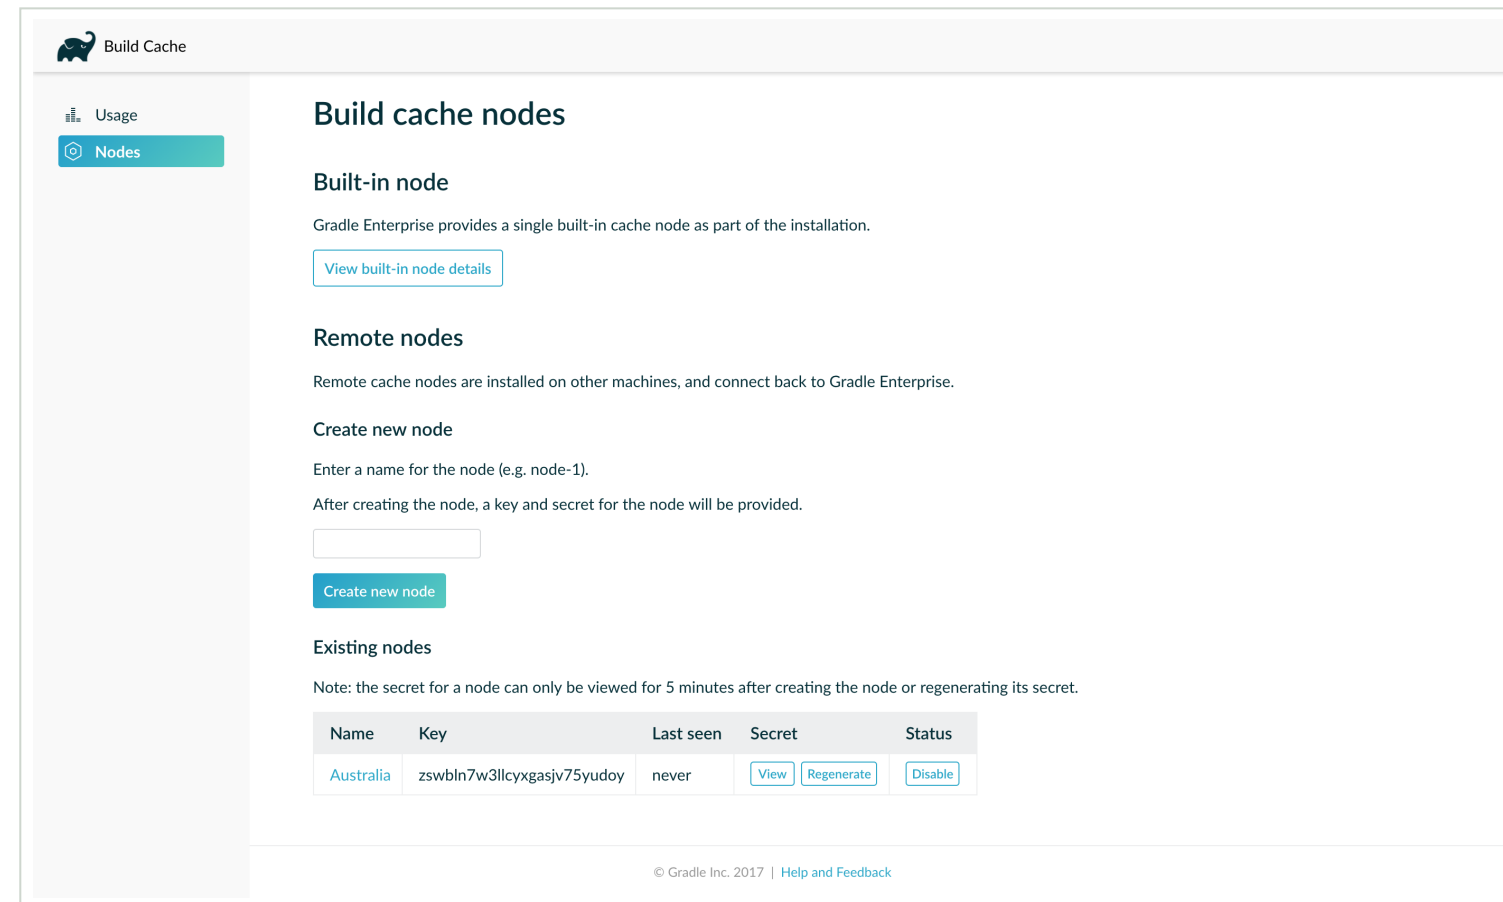
```
docker pull gradle/build-cache-node
```

Owner

 gradle

Connecting to Gradle Enterprise

- Optionally **register** with Gradle Enterprise for centralized management
- **Replication capabilities** for geographically distributed teams



The screenshot displays the 'Build Cache' interface in Gradle Enterprise. On the left, a sidebar contains 'Usage' and 'Nodes' (selected). The main content area is titled 'Build cache nodes' and includes sections for 'Built-in node', 'Remote nodes', and 'Create new node'. The 'Create new node' section has a text input and a 'Create new node' button. Below, the 'Existing nodes' section shows a table with one node named 'Australia'.

Build Cache

Build cache nodes

Built-in node

Gradle Enterprise provides a single built-in cache node as part of the installation.

[View built-in node details](#)

Remote nodes

Remote cache nodes are installed on other machines, and connect back to Gradle Enterprise.

Create new node

Enter a name for the node (e.g. node-1).

After creating the node, a key and secret for the node will be provided.

[Create new node](#)

Existing nodes

Note: the secret for a node can only be viewed for 5 minutes after creating the node or regenerating its secret.

Name	Key	Last seen	Secret	Status
Australia	zswbln7w3llcyxgasjv75yudoy	never	View Regenerate	Disable

© Gradle Inc. 2017 | [Help and Feedback](#)

Wrap up

Documentation and resources

- gradle.com/enterprise/resources
- docs.gradle.org/current/userguide/build_cache.html
- guides.gradle.org/using-build-cache

Video playlist

To review concepts and learn more about build cache and distributed testing, check out our brand new video playlist called **Faster Feedback Cycles**

tv.gradle.com/build-faster

Try Gradle Enterprise for free

- Free 30-day trial
 - hosted by Gradle zero installation, ready-to-go
 - on your infrastructure quick setup, maximum control
- Technical support included

<https://gradle.com/enterprise/trial>



Thank you